


## Aula 6:


Redes Convolucionais – Parte II - VGG,  
GoogLeNet, Inception Module, ResNet,  
Highways networks.....

[www.deeplearningbrasil.com.br](http://www.deeplearningbrasil.com.br)




# Relembrando...

## Desafio ImageNet: top-5 error




# Arquitetura antes do conceito de Deep Learning

---




# Arquitectura convolucional

---



# Arquitectura convolucional

---



# Arquitetura convolucional da AlexNet


---

- Relembrando as camadas:
  - Entrada  $224 \times 224 \times 3$
  - Primeiro filtro  $11 \times 11 \times 3$
  - Filtros convolucionais menores a medida que aumenta a profundidade
  - Ex.:  $5 \times 5$ ,  $3 \times 3$
- O que acontece se usarmos apenas filtros pequenos? ( $3 \times 3$ , por exemplo?)

# VGGNet: ILSVRC 2014 2<sup>nd</sup> Lugar

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 <b>LRN</b>	conv3-64 <b>conv3-64</b>	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 <b>conv3-128</b>	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 <b>conv1-256</b>	conv3-256 conv3-256 <b>conv3-256</b>	conv3-256 conv3-256 conv3-256 <b>conv3-256</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

- Sequence of deeper networks trained progressively
- Campos receptivos substituídos por filtros 3x3 (com ReLU entre eles)



- Uma camada 7x7 com C mapas precisa de  $49C^2$  pesos, camadas 3x3 precisam de  $27C^2$  pesos


Table 2: Number of parameters (in millions).

Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

K. Simonyan and A. Zisserman, [Very Deep Convolutional Networks for Large-Scale Image Recognition](#), ICLR 2015

# VGGNet: ILSVRC 2014 2<sup>nd</sup> Lugar

---






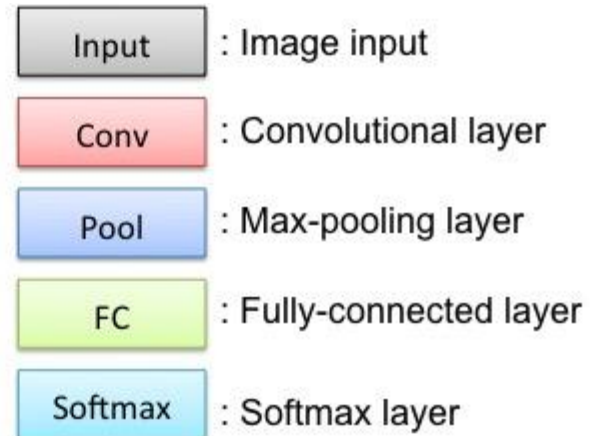

# Características da VGGNet

---

## AlexNet



## VGGNet



# Codificando...

---

```
ImageNet classification with Python and Keras
1  # import the necessary packages
2  from keras.preprocessing import image as image_utils
3  from imagenet_utils import decode_predictions
4  from imagenet_utils import preprocess_input
5  from vgg16 import VGG16
6  import numpy as np
7  import argparse
8  import cv2
9
10 # construct the argument parse and parse the arguments
11 ap = argparse.ArgumentParser()
12 ap.add_argument("-i", "--image", required=True,
13                 help="path to the input image")
14 args = vars(ap.parse_args())
15
16 # load the original image via OpenCV so we can draw on it and display
17 # it to our screen later
18 orig = cv2.imread(args["image"])
```

Exemplo obtido de:

<http://www.pyimagesearch.com/2016/08/10/imagenet-classification-with-python-and-keras/>

# Codificando...

---

ImageNet classification with Python and Keras

Python

```
20 # load the input image using the Keras helper utility while ensuring
21 # that the image is resized to 224x224 pixels, the required input
22 # dimensions for the network -- then convert the PIL image to a
23 # NumPy array
24 print("[INFO] loading and preprocessing image...")
25 image = image_utils.load_img(args["image"], target_size=(224, 224))
26 image = image_utils.img_to_array(image)
```

# Codificando...

---

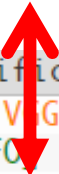
ImageNet classification with Python and Keras

```
28 # our image is now represented by a NumPy array of shape (3, 224, 224),  
29 # but we need to expand the dimensions to be (1, 3, 224, 224) so we can  
30 # pass it through the network -- we'll also preprocess the image by  
31 # subtracting the mean RGB pixel intensity from the ImageNet dataset  
32 image = np.expand_dims(image, axis=0)  
33 image = preprocess_input(image)
```

# Codificando...

---

VGG16, VGG19 ou ResNet50



```
ImageNet classification with Python and Keras
35 # load the VGG16 network
36 print("[INFO] loading network...")
37 model = VGG16(weights="imagenet")
38
39 # classify the image
40 print("[INFO] classifying image...")
41 preds = model.predict(image)
42 (inID, label) = decode_predictions(preds)[0]
43
44 # display the predictions to our screen
45 print("ImageNet ID: {}, Label: {}".format(inID, label))
46 cv2.putText(orig, "Label: {}".format(label), (10, 30),
47            cv2.FONT_HERSHEY_SIMPLEX, 0.9, (0, 255, 0), 2)
48 cv2.imshow("Classification", orig)
49 cv2.waitKey(0)
```

```
$ python test_imagenet.py --image images/dog_beagle.png
```

---

**Caramba...  
esse curso é demais...  
Não preciso de mais nada...**

**Calma...**



# GoogLeNet: ILSVRC 2014 winner

---




<http://knowyourmeme.com/memes/we-need-to-go-deeper>

C. Szegedy et al., [Going deeper with convolutions](#), CVPR 2015

---


**Em imagens, correlações tender a ser local**






---

**Ideia: cobrir clusters locais com convoluções 1x1**




---

## Correlações menos espalhadas




---

## Correlações menos espalhadas




---


## Correlações menos espalhadas



---

## Correlações menos espalhadas

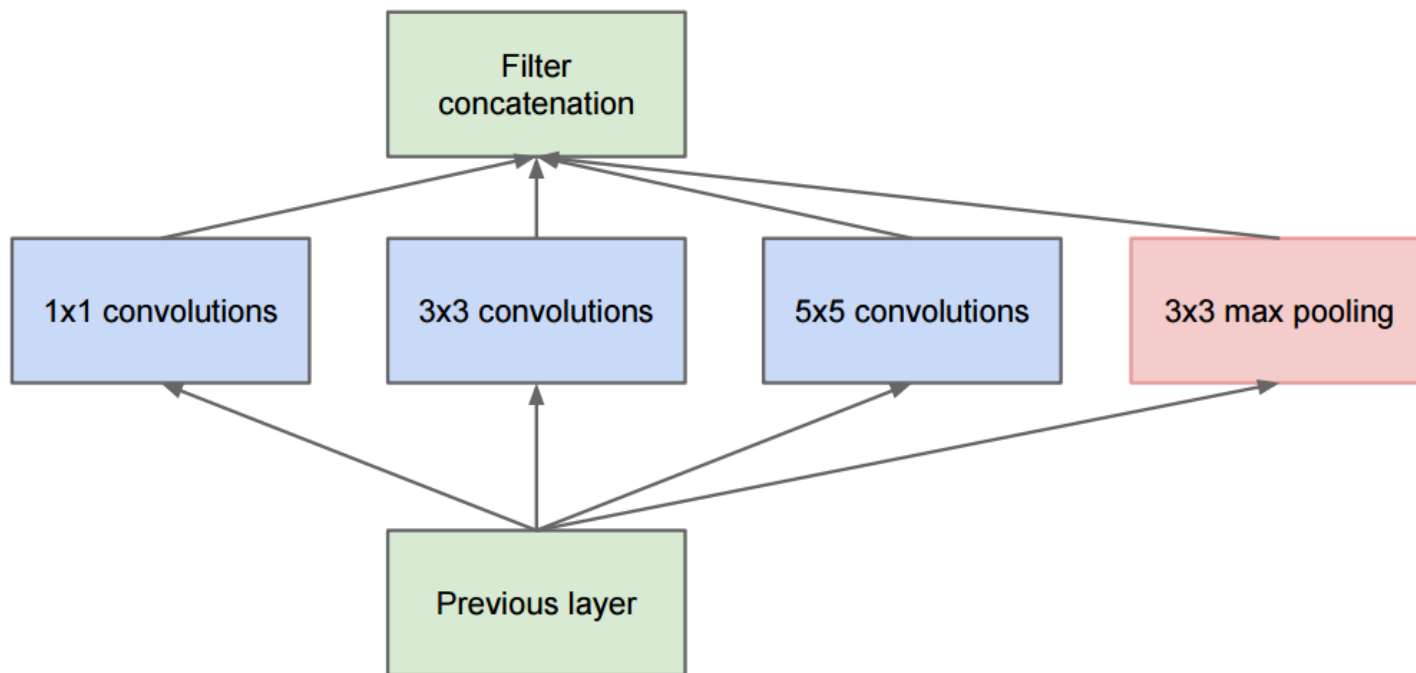




# GoogLeNet

---

- Inception Module v1
  - Filtros convolucionais em paralelo com campos receptivos diferentes
  - Usa filtros 1x1 para redução de dimensionalidade antes de operações de “alto custo” antes das convoluções maiores



# GoogLeNet

---

## Convolução 1x1


- Reduz o esforço computacional antes de filtros 3x3 e 5x5
- Inclui o uso de RELU
- Exemplo:
- Entrada 256 mapas de filtros 56x56
- Inserção de camada 1x1x64
- 56x56x256 -> 56x56x64 (redução de 4x)

<http://iamaaditya.github.io/2016/03/one-by-one-convolution/>



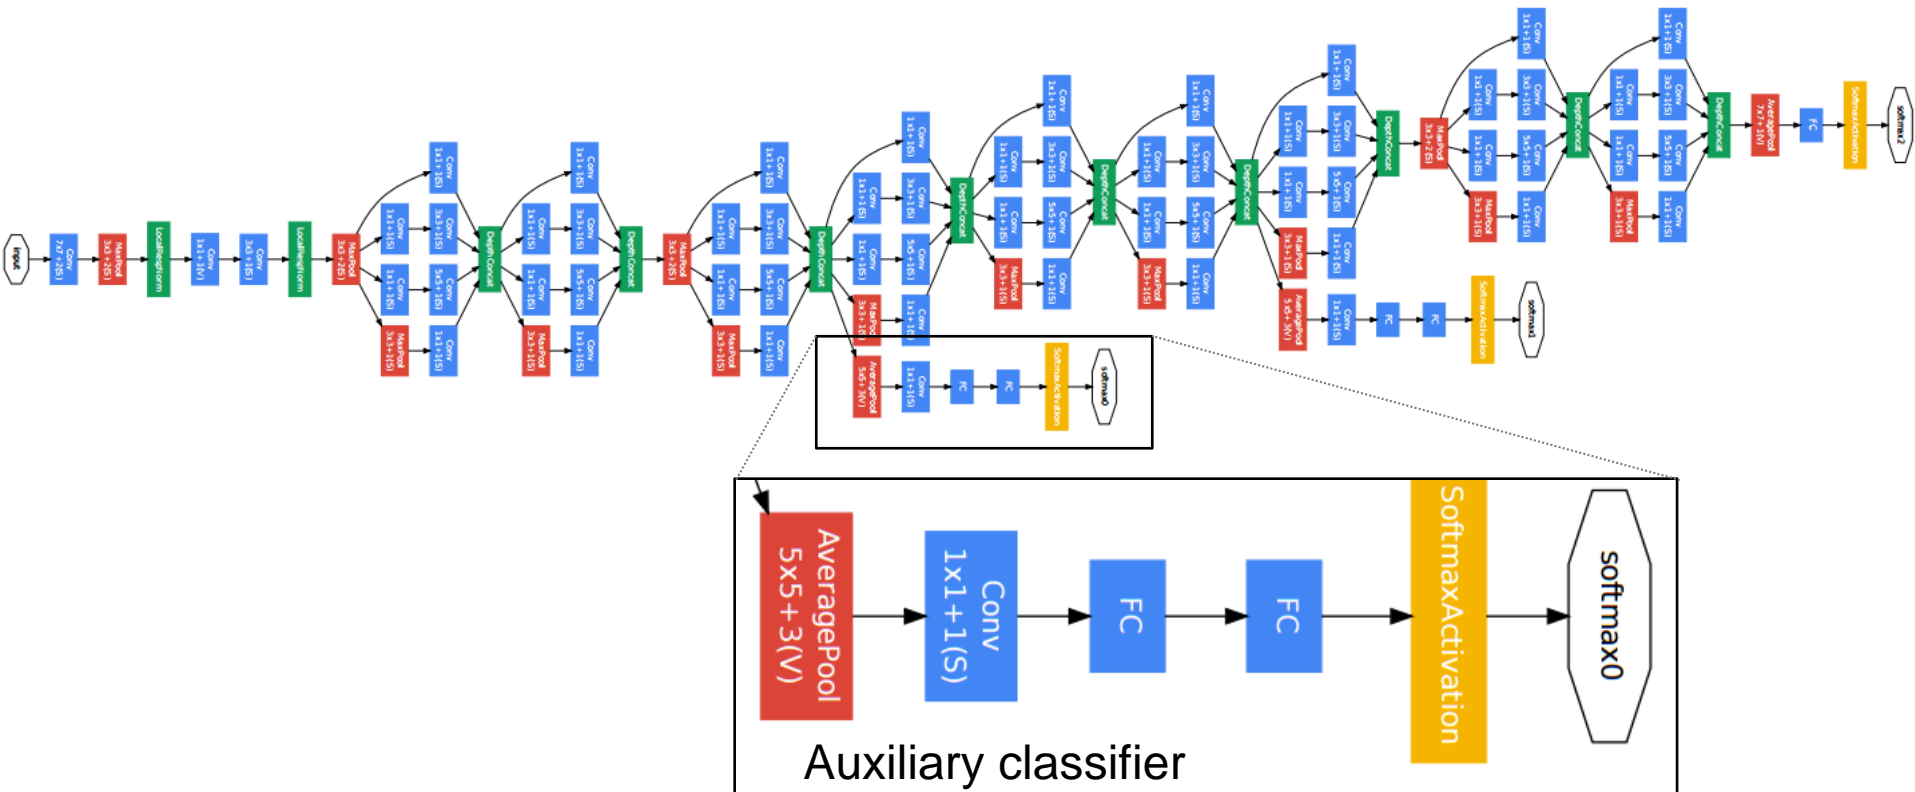
# GoogLeNet

---





# GoogLeNet



C. Szegedy et al., [Going deeper with convolutions](#), CVPR 2015

# GoogLeNet

---

An alternative view:

type	patch size/ stride	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj	params	ops
convolution	7×7/2	112×112×64	1							2.7K	34M
max pool	3×3/2	56×56×64	0								
convolution	3×3/1	56×56×192	2		64	192				112K	360M
max pool	3×3/2	28×28×192	0								
inception (3a)		28×28×256	2	64	96	128	16	32	32	159K	128M
inception (3b)		28×28×480	2	128	128	192	32	96	64	380K	304M
max pool	3×3/2	14×14×480	0								
inception (4a)		14×14×512	2	192	96	208	16	48	64	364K	73M
inception (4b)		14×14×512	2	160	112	224	24	64	64	437K	88M
inception (4c)		14×14×512	2	128	128	256	24	64	64	463K	100M
inception (4d)		14×14×528	2	112	144	288	32	64	64	580K	119M
inception (4e)		14×14×832	2	256	160	320	32	128	128	840K	170M
max pool	3×3/2	7×7×832	0								
inception (5a)		7×7×832	2	256	160	320	32	128	128	1072K	54M
inception (5b)		7×7×1024	2	384	192	384	48	128	128	1388K	71M
avg pool	7×7/1	1×1×1024	0								
dropout (40%)		1×1×1024	0								
linear		1×1×1000	1							1000K	1M
softmax		1×1×1000	0								

C. Szegedy et al., [Going deeper with convolutions](#), CVPR 2015

## Rethinking the Inception Architecture for Computer Vision

[Christian Szegedy](#), [Vincent Vanhoucke](#), [Sergey Ioffe](#), [Jonathon Shlens](#), [Zbigniew Wojna](#)

*(Submitted on 2 Dec 2015 (v1), last revised 11 Dec 2015 (this version, v3))*

Convolutional networks are at the core of most state-of-the-art computer vision solutions for a wide variety of tasks. Since 2014 very deep convolutional networks started to become mainstream, yielding substantial gains in various benchmarks. Although increased model size and computational cost tend to translate to immediate quality gains for most tasks (as long as enough labeled data is provided for training), computational efficiency and low parameter count are still enabling factors for various use cases such as mobile vision and big-data scenarios. Here we explore ways to scale up networks in ways that aim at utilizing the added computation as efficiently as possible by suitably factorized convolutions and aggressive regularization. We benchmark our methods on the ILSVRC 2012 classification challenge validation set demonstrate substantial gains over the state of the art: 21.2% top-1 and 5.6% top-5 error for single frame evaluation using a network with a computational cost of 5 billion multiply-adds per inference and with using less than 25 million parameters. With an ensemble of 4 models and multi-crop evaluation, we report 3.5% top-5 error on the validation set (3.6% error on the test set) and 17.3% top-1 error on the validation set.

Subject: Computer Vision and Pattern Recognition (cs.CV)

---

## Princípios de design:

- 1) Avoid Representational Bottlenecks,  
Especially Early in the Network
- 2) Representações dimensionais mais altas  
são mais fáceis de processar localmente
- 3) A agregação espacial pode ser feita em  
dimensões mais baixas sem perda  
Ex.: Antes de 3x3, redução de dimensionalidade
- 4) Equilíbrio entre profundidade e largura


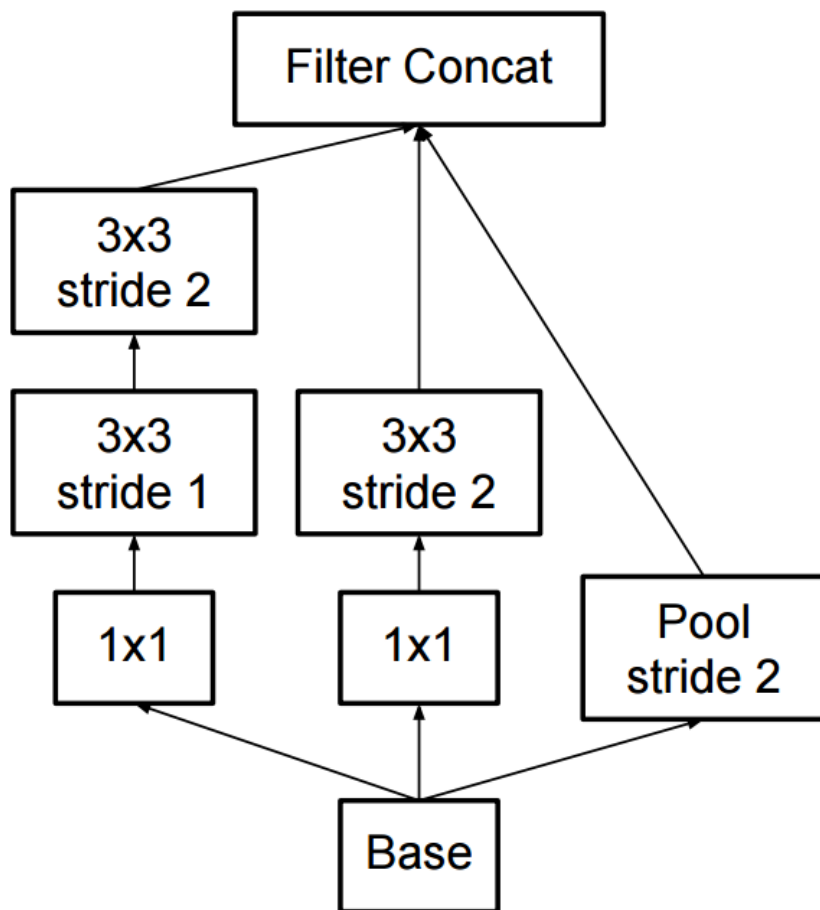
---

Ponto chave: Fatorização das convoluções

# Inception v2

---

Inception v2 substituiu as convoluções 5x5 por dois filtros 3x3 sucessivos:




**Figura 5**



# Inception v2

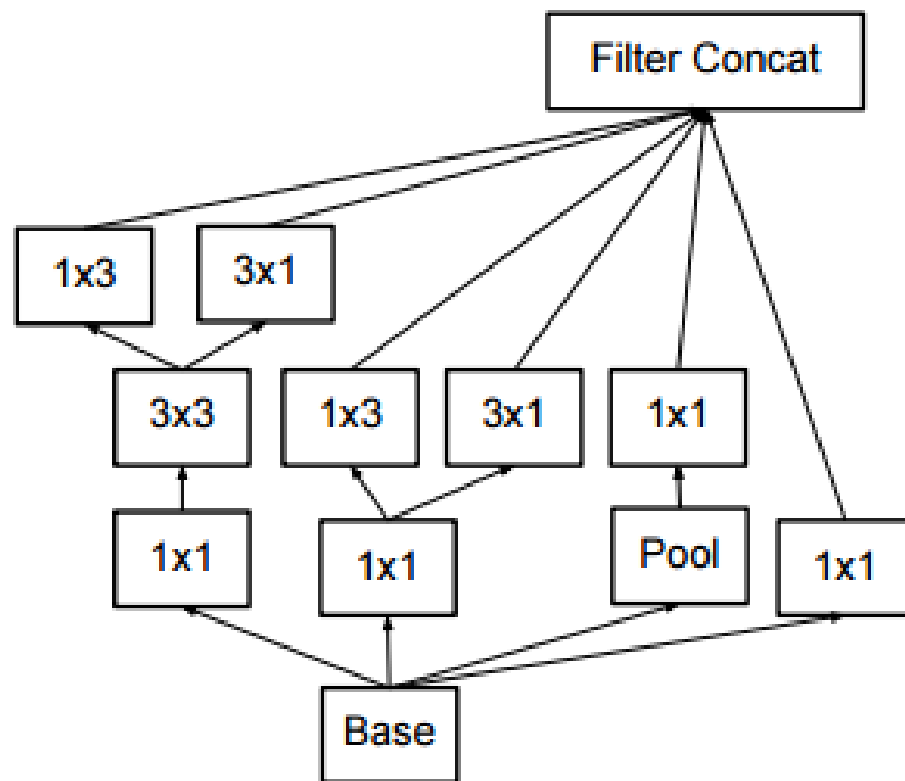
- Mais variantes dos módulos com fatorização mais agressiva de filtros



**Figura 6**

# Inception v2

- Mais variantes dos módulos com fatorização mais agressiva de filtros

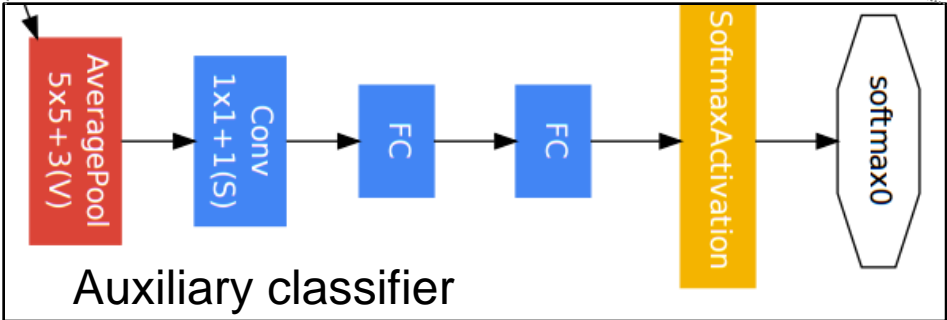


**Figura 7**

# Inception v2/v3

---

type	patch size/stride or remarks	input size
conv	$3 \times 3 / 2$	$299 \times 299 \times 3$
conv	$3 \times 3 / 1$	$149 \times 149 \times 32$
conv padded	$3 \times 3 / 1$	$147 \times 147 \times 32$
pool	$3 \times 3 / 2$	$147 \times 147 \times 64$
conv	$3 \times 3 / 1$	$73 \times 73 \times 64$
conv	$3 \times 3 / 2$	$71 \times 71 \times 80$
conv	$3 \times 3 / 1$	$35 \times 35 \times 192$
3×Inception	As in figure 5	$35 \times 35 \times 288$
5×Inception	As in figure 6	$17 \times 17 \times 768$
2×Inception	As in figure 7	$8 \times 8 \times 1280$
pool	$8 \times 8$	$8 \times 8 \times 2048$
linear	logits	$1 \times 1 \times 2048$
softmax	classifier	$1 \times 1 \times 1000$



C. Szegedy et al., [Going deeper with convolutions](#), CVPR 2015

**I WAS WINNING  
IMAGENET**



**UNTIL A  
DEEPER MODEL  
CAME ALONG**

[Source \(?\)](#)

# ResNet: ILSVRC 2015 winner

---

## Revolution of Depth

AlexNet, 8 layers  
(ILSVRC 2012)



VGG, 19 layers  
(ILSVRC 2014)




ResNet, 152 layers  
(ILSVRC 2015)

Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun, [Deep Residual Learning for Image Recognition](#), CVPR 2016

# Problema:

---

Até o momento: empilhar convoluções 3x3  
Funciona bem até 20~30 camadas



# Problema

---

Para problemas mais complexos

- Campos receptive (convoluções)  $\uparrow$
- Não linearidade  $\uparrow$

Entretanto, quanto maior a “profundidade”:

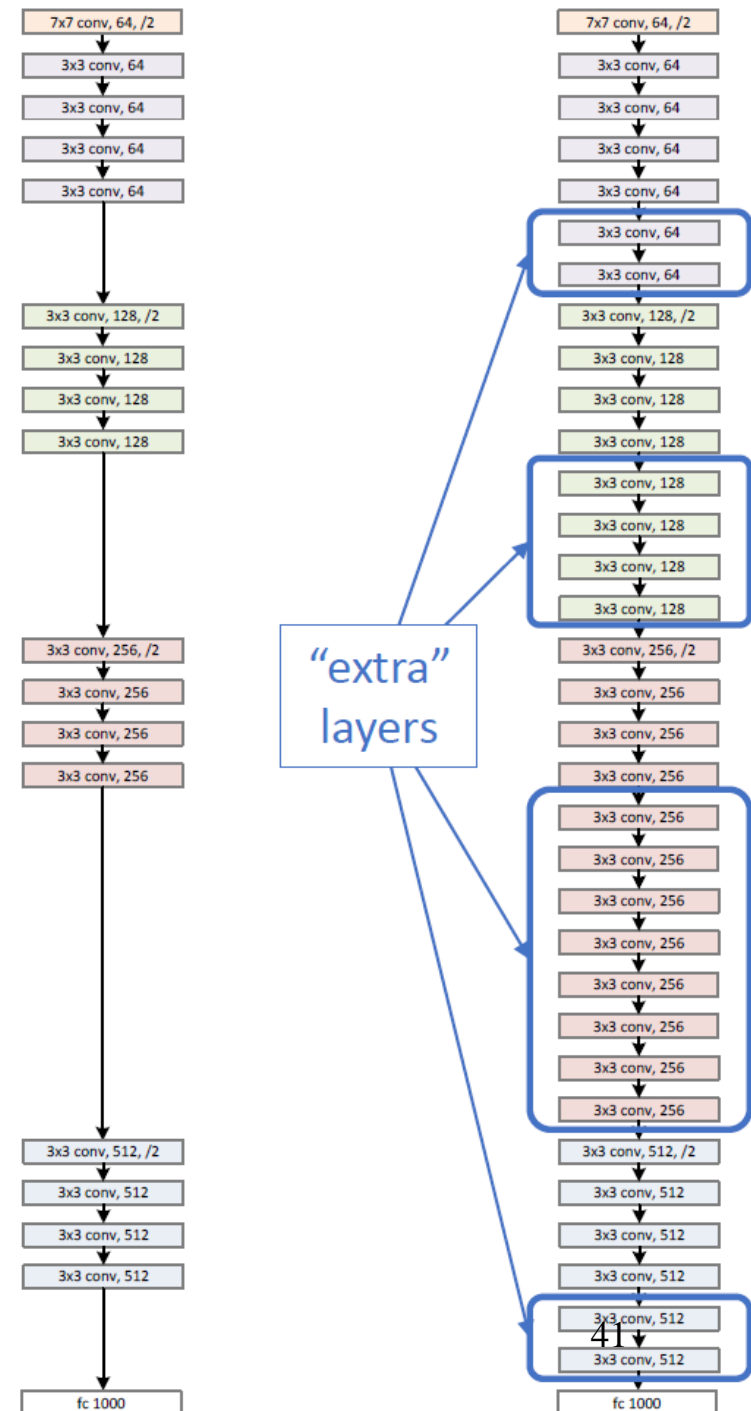
- vanishing/exploding gradients problem



# ResNet

## Solução Naïve

- Se as camadas extras forem um mapeamento de identidade, os erros de treinamento não aumentarão





# ResNet

---

- Módulo residual
  - Introduz conexões de “atalho”

**Ideia do neurônio clássico**



$$x^k = f((W^k * x^{k-1}) + b^k).$$

# ResNet

---

- Módulo residual
  - Introduz conexões de “atalho”


**Ideia do neurônio clássico**


$$x^k = f((W^k * x^{k-1}) + b^k).$$

**Módulo residual**



$$y^k = F(x^k, W^k) + x^k,$$



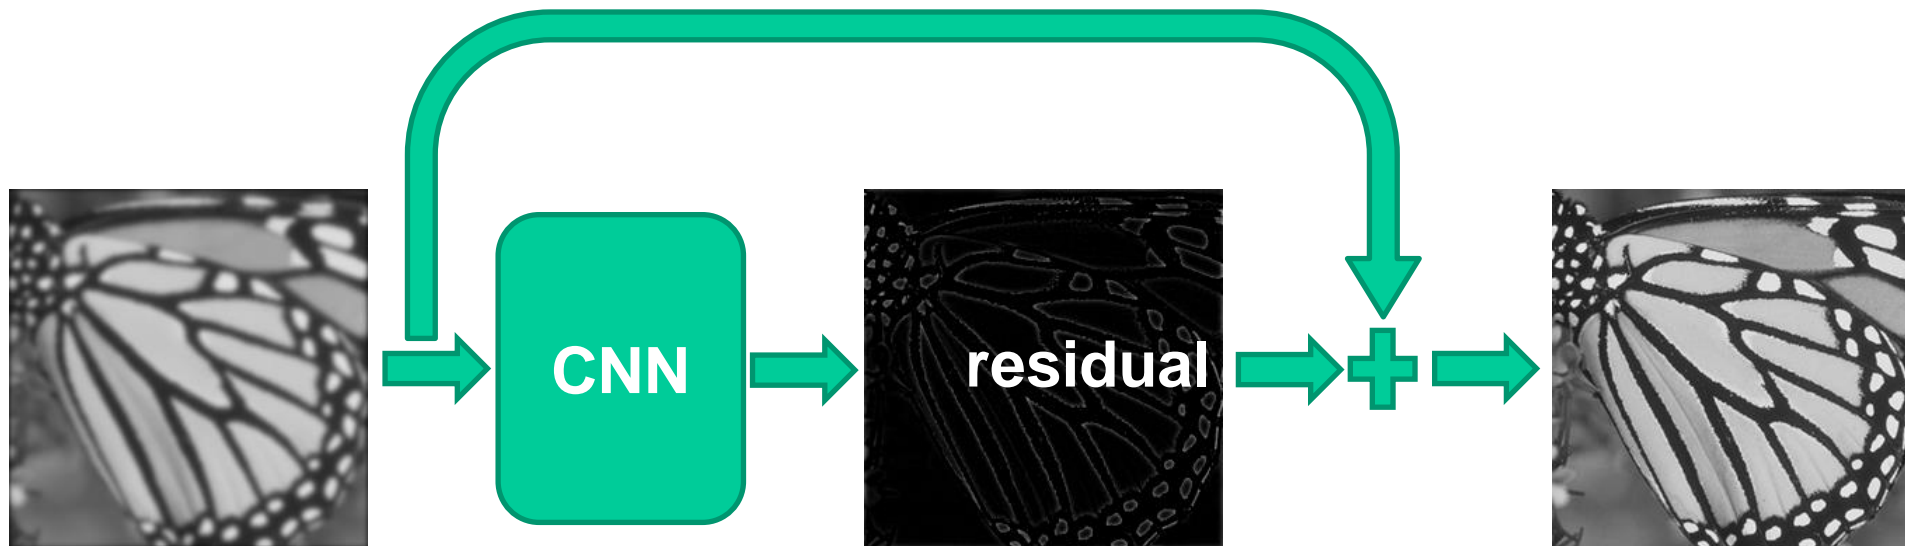
# ResNet

---

- Cada pilha de camadas se encaixe diretamente no mapeamento subjacente desejado
- Explicitamente permitimos que essas camadas se encaixem em um mapeamento residual.
- Hipótese de que é mais fácil otimizar o mapeamento residual do que otimizar o mapeamento original.

# ResNet

---




# ResNet


---

## Bloco Residual

- Simples
- Livre de parâmetro




**Bloco “naïve”**



**Bloco “bottleneck”**  
(ResNet-50/101/152)

# ResNet

## Módulo “bottleneck”



- Com convoluções  $3 \times 3$  e 256 mapas de entrada:  
 $256 \times 256 \times 3 \times 3 \sim 600K$
- Com convoluções  $1 \times 1$  para reduzir de 256 para 64 mapas, seguidos por convoluções  $3 \times 3$ , seguidos por convoluções  $1 \times 1$  para expandir para 256 mapas:  
 $256 \times 64 \times 1 \times 1 \sim 16K$   
 $64 \times 64 \times 3 \times 3 \sim 36K$   
 $64 \times 256 \times 1 \times 1 \sim 16K$   
Total:  $\sim 70K$



# ResNet

---

## Alguns tipos de “atalhos”

- Identidade

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{W_i\}) + \mathbf{x}.$$

- Projeção

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{W_i\}) + W_s \mathbf{x}.$$

# ResNet

## Arquitetura para ImageNet:

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
conv2_x	56×56	3×3 max pool, stride 2				
		$\begin{bmatrix} 3\times 3, 64 \\ 3\times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 64 \\ 3\times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3\times 3, 128 \\ 3\times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 128 \\ 3\times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3\times 3, 256 \\ 3\times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 256 \\ 3\times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3\times 3, 512 \\ 3\times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 512 \\ 3\times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun, [Deep Residual Learning for Image Recognition](#), CVPR 2016 (Best Paper)

# Inception v4

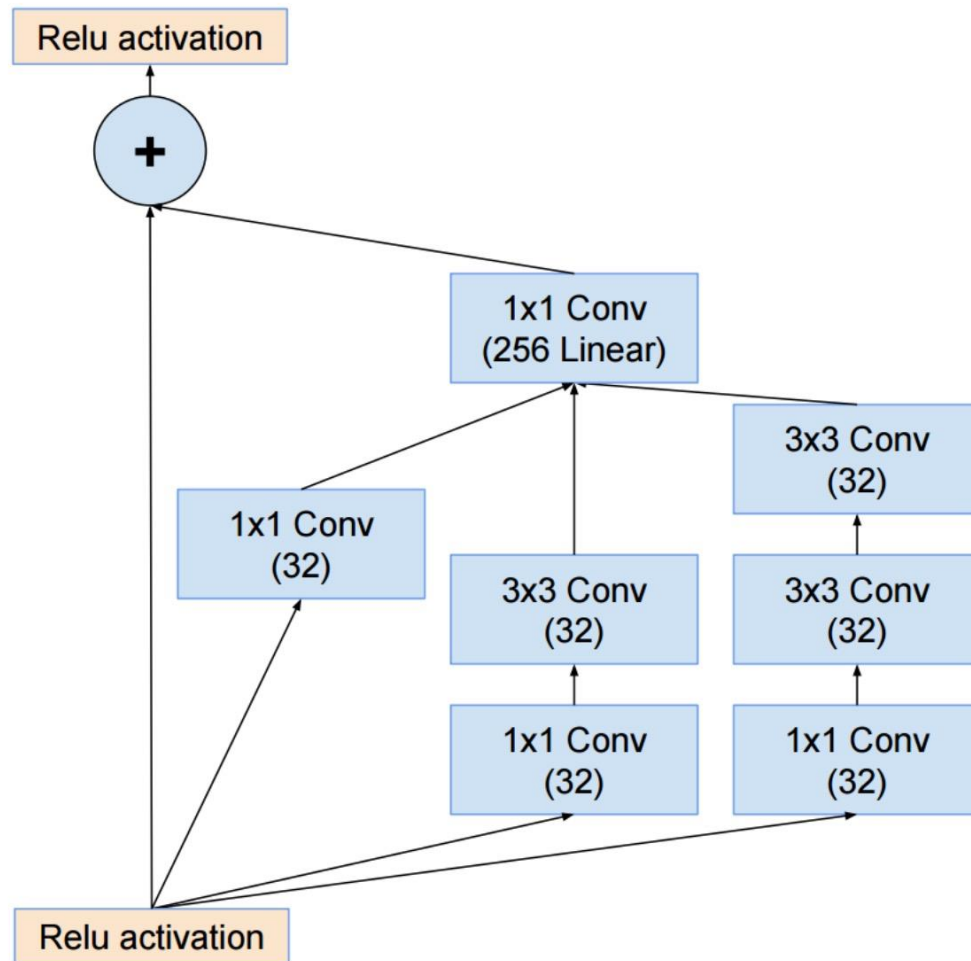
---

**Duas “redes” poderosas:**


- Inception
- Residual

# Inception v4


---



C. Szegedy et al., [Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning](#), arXiv 2016



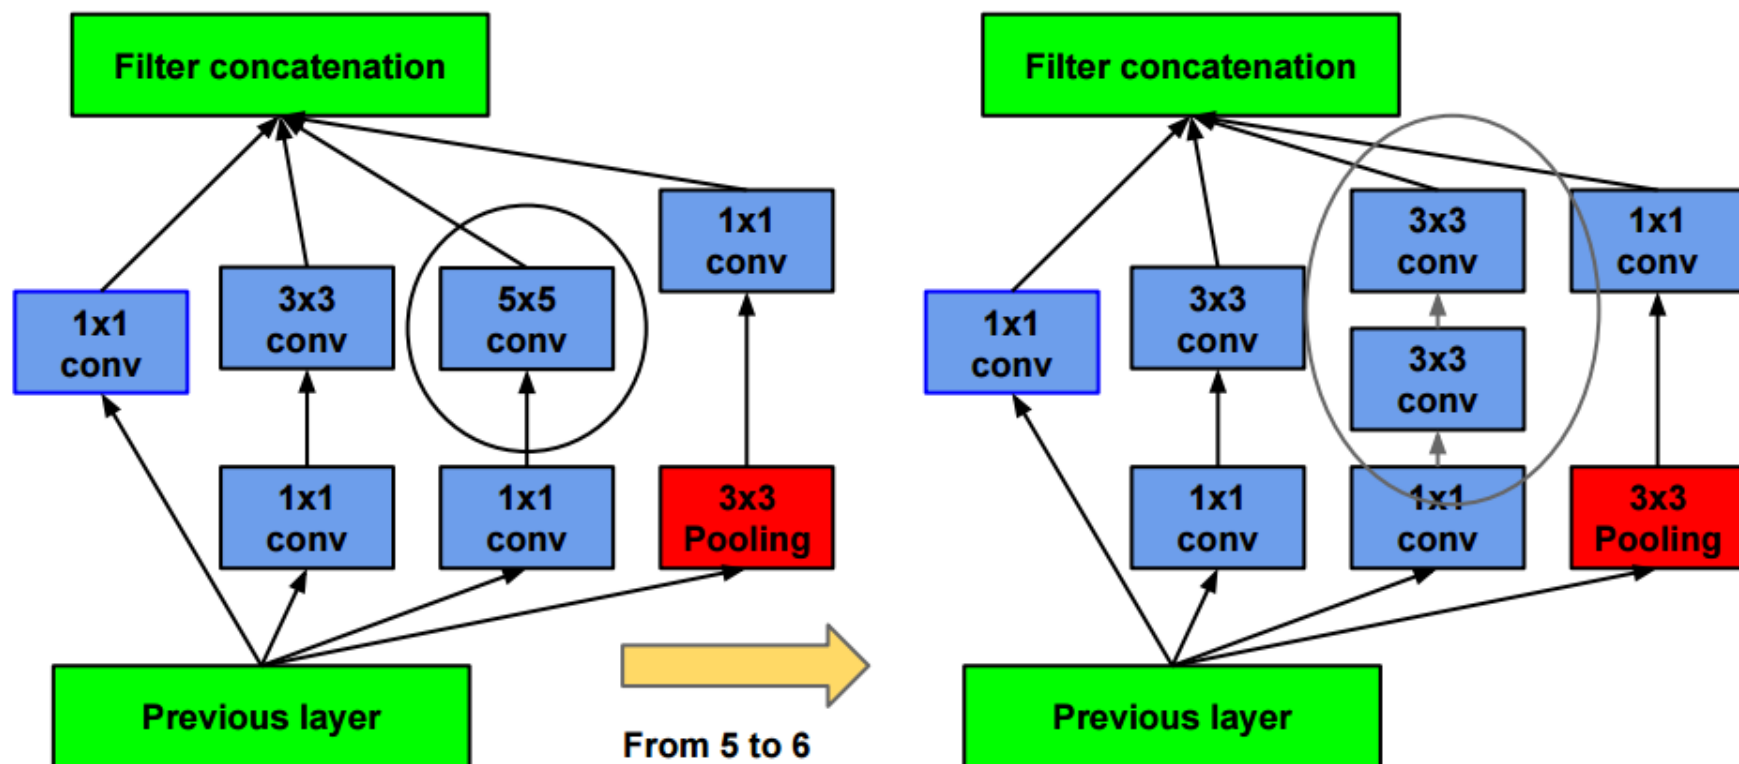
<sup>1</sup>Inception 5 (GoogLeNet)



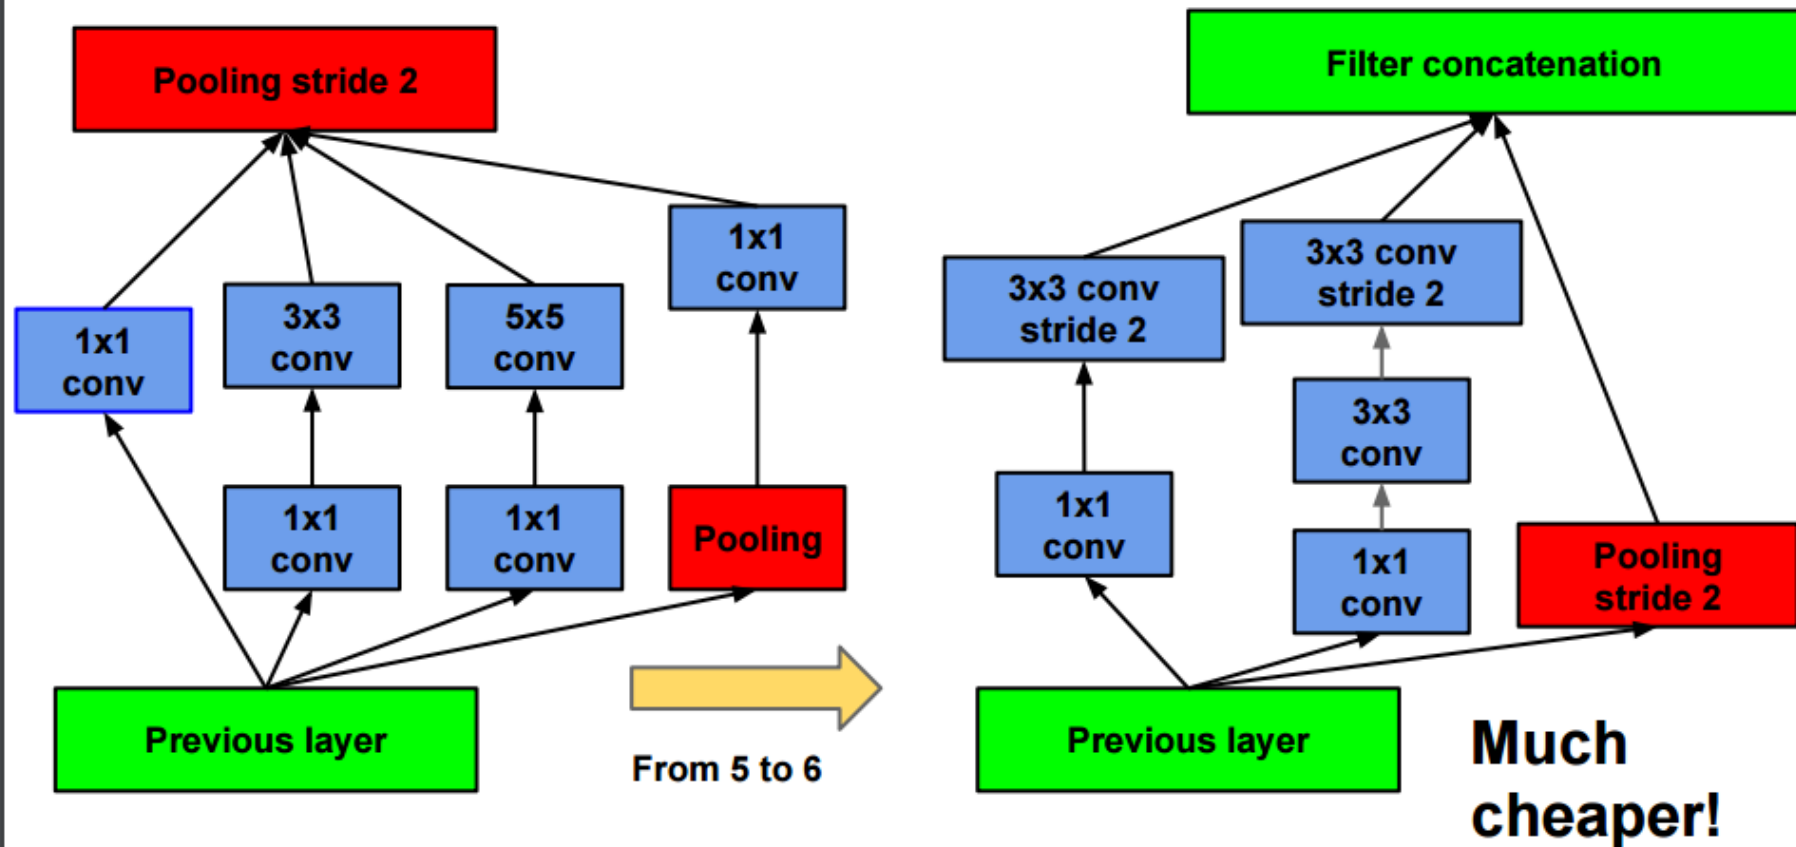
Inception 7a

<sup>1</sup>Going Deeper with Convolutions, [C. Szegedy et al, CVPR 2015]

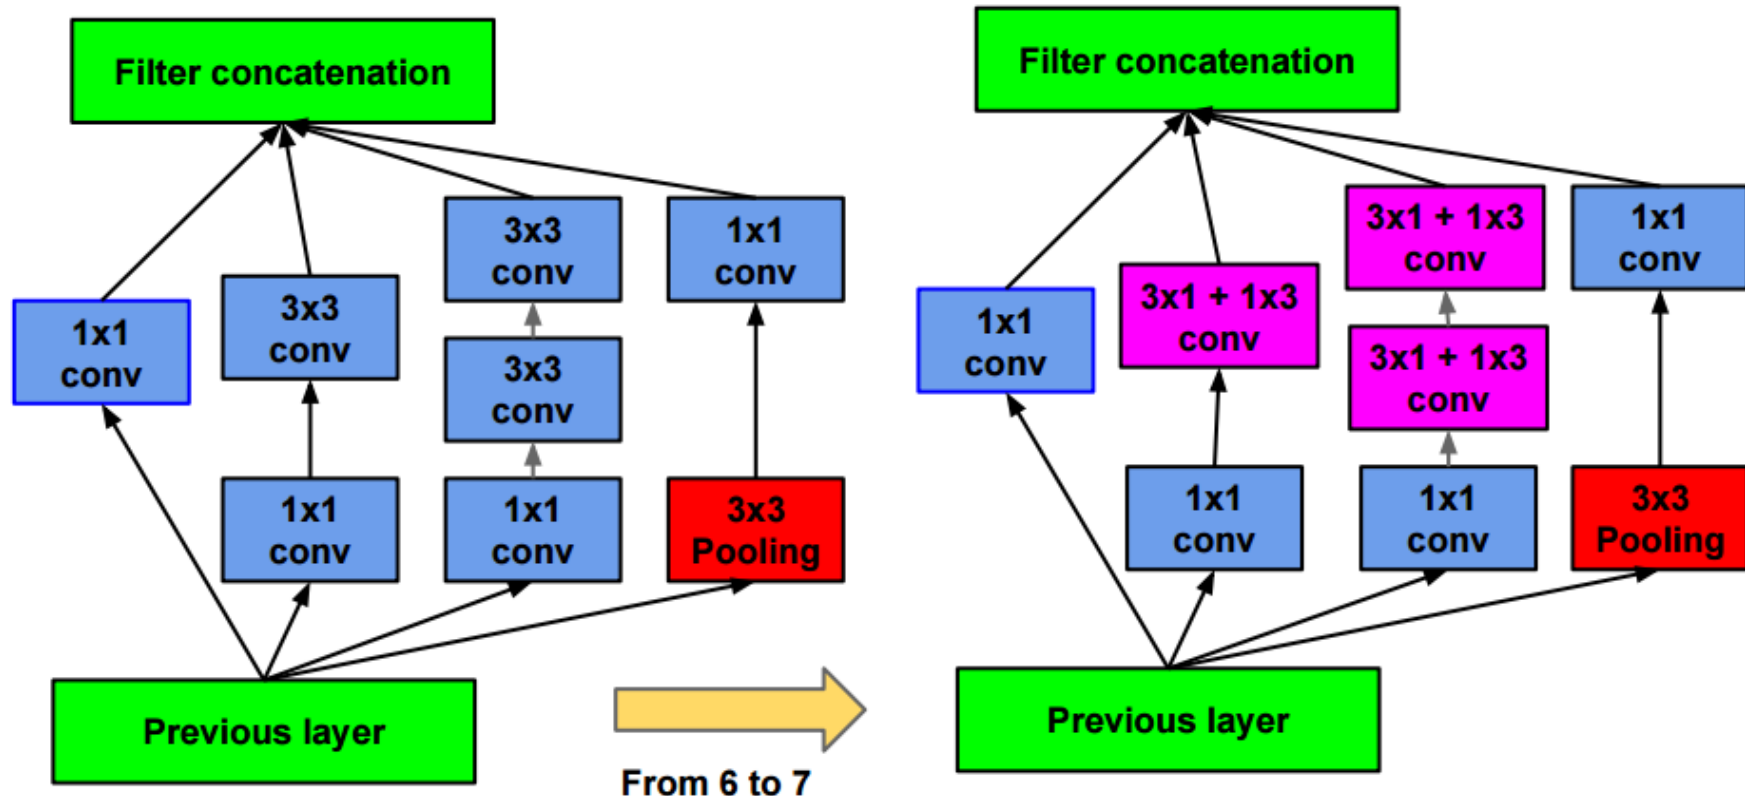
# Structural changes from Inception 5 to 6



# Grid size reduction Inception 5 vs 6



# Structural changes from Inception 6 to 7






# Summary: ILSVRC 2012-2015

---

Team	Year	Place	Error (top-5)	External data
SuperVision – Toronto (AlexNet, 7 layers)	2012	-	16.4%	no
SuperVision	2012	1st	15.3%	ImageNet 22k
Clarifai – NYU (7 layers)	2013	-	11.7%	no
Clarifai	2013	1st	11.2%	ImageNet 22k
VGG – Oxford (16 layers)	2014	2nd	7.32%	no
GoogLeNet (19 layers)	2014	1st	6.67%	no
ResNet (152 layers)	2015	1st	3.57%	
Human expert*			5.1%	

<http://karpathy.github.io/2014/09/02/what-i-learned-from-competing-against-a-convnet-on-imagenet/>

# Accuracy vs. efficiency




LeNet



AlexNet



VGG



GoogLeNet



Inception V3



Inception BN




```

graph BT
    A[Convolution 5x5, 32] -- "32x256x256" --> B[Activation tanh]
    B -- "32x256x256" --> C[Pooling max, 2x2]
    C -- "32x140x140" --> D[Convolution 5x5, 256]
    D -- "32x140x140" --> E[Activation tanh]
    E -- "32x140x140" --> F[Pooling max, 2x2]
    F -- "32x72x72" --> G[Flatten]
    G -- "29200" --> H[Fully Connected 500]
    H -- "500" --> I[Activation tanh]
    I -- "500" --> J[Fully Connected 50]
    J -- "50" --> K[Softmax/Output]
  
```

[illegible]

Fully Connected 2

[illegible]

# Highway Networks

---

Para uma camada densa “convencional”

$$y = H(x, W_H)$$

```
def dense(x, input_size, output_size, activation):  
    W = tf.Variable(tf.truncated_normal([input_size, output_size], stddev=0.1), name="weight")  
    b = tf.Variable(tf.constant(0.1, shape=[output_size]), name="bias")  
    y = activation(tf.matmul(x, W) + b)  
    return y
```

# Highway Networks

---

A camada tem dois gateways que controlam o fluxo de informações. O gate que controla o quanto da ativação passa para a próxima camada... e o gate que controla o quanto da entrada sem modificação passa "direto" para a próxima camada.

$$\mathbf{y} = H(\mathbf{x}, \mathbf{W}_H) \cdot T(\mathbf{x}, \mathbf{W}_T) + \mathbf{x} \cdot C(\mathbf{x}, \mathbf{W}_C)$$

- Um conjunto extra de pesos e bias para aprender os “gates”
- A operação de transformação ( $T$ ).
- “carry gate” ( $C$  equivalente a  $1 - T$ ).

# Highway Networks

---

Para saber mais...

---

## Training Very Deep Networks


---

Rupesh Kumar Srivastava   Klaus Greff   Jürgen Schmidhuber

The Swiss AI Lab IDSIA / USI / SUPSI  
{rupesh, klaus, juergen}@idsia.ch

# Densenet


---







# Densenet

---




# Densenet

---



# Densenet

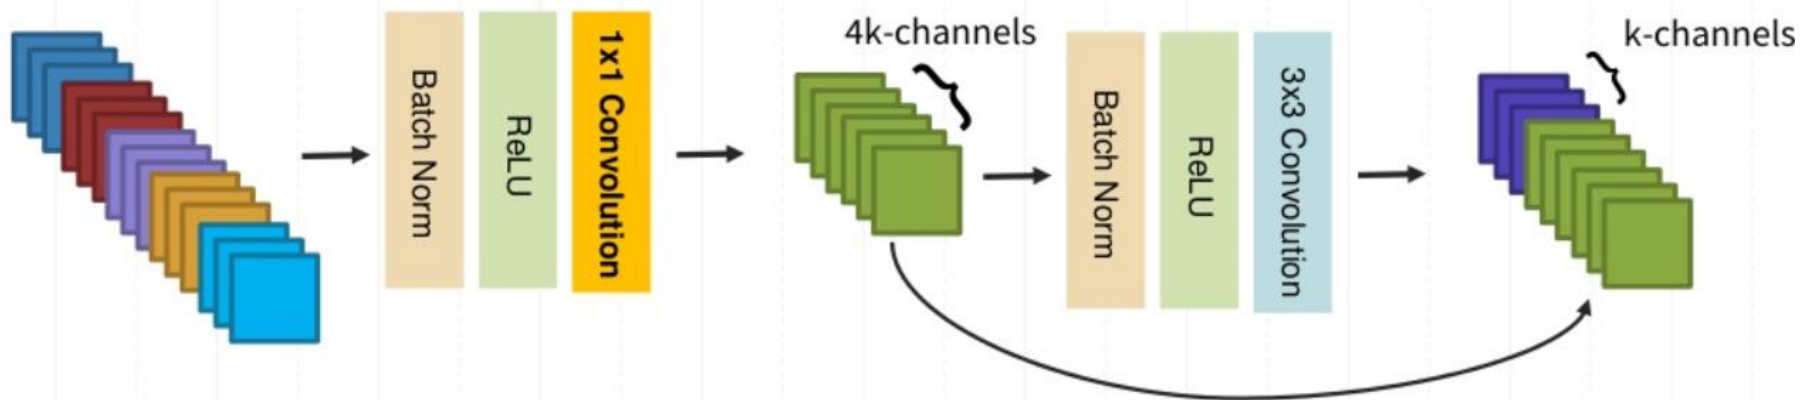
---



# Densenet

---

- Reduce the number of channels to  $4k$  before composite layer by **1x1 convolution**.
- It's optional.




---


# Transfer Learning

# Transfer Learning

---



**Gato**



**Cachorro**



**elephant**



**tigre**

**Domínio similar, tarefas  
diferentes**

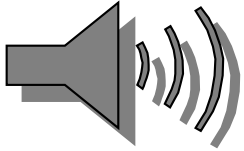

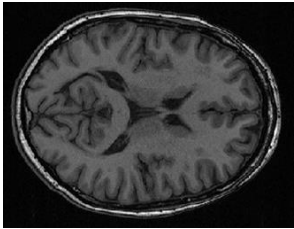





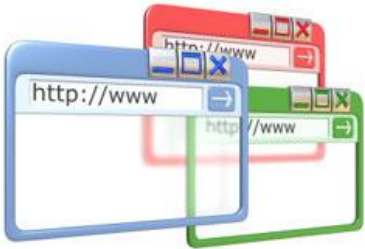


**dog**

**cat**

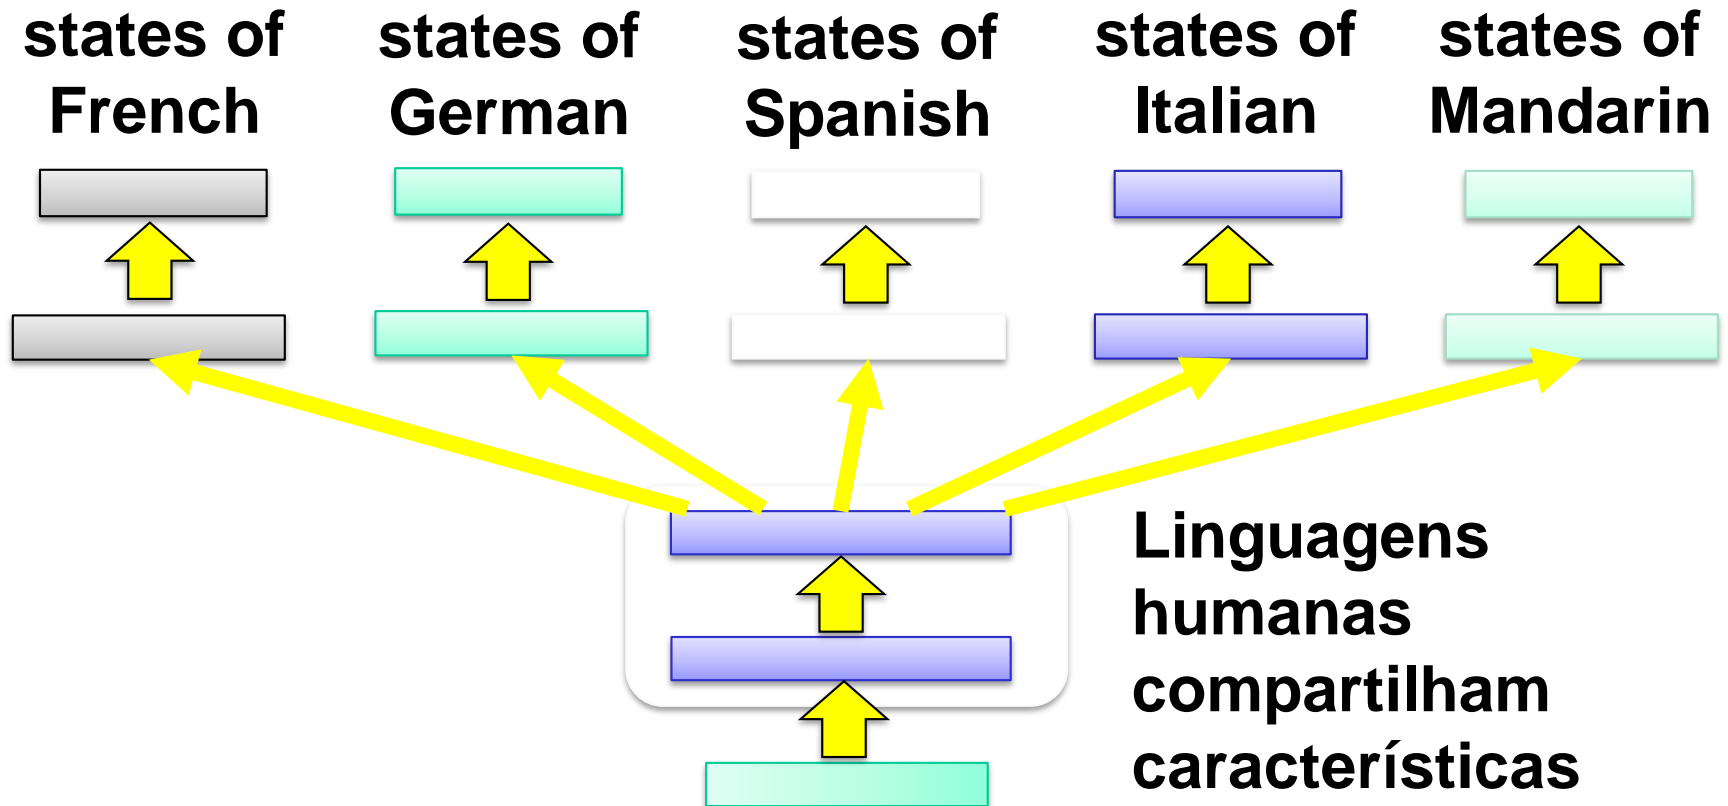
**Domínio diferente,  
mesma tarefa**

# Transfer Learning

Tarefa		Dados
Speech Recognition	 Taiwanese	 English Chinese .....
Image Recognition	 Medical Images	   
Text Analysis	 Specific domain	 Webpages

# Transfer Learning

---





**Similar idea in translation:** Daxiang Dong, Hua Wu, Wei He, Dianhai Yu and Haifeng Wang, "Multi-task learning for multiple language translation.", ACL 2015



# Transfer Learning Model Fine-tuning

---

## Tarefa

- Source data:  $(x^s, y^s)$   **Muito**
- Target data:  $(x^t, y^t)$   **Pouco**

## Exemplos: reconhecimento de fala e melanoma

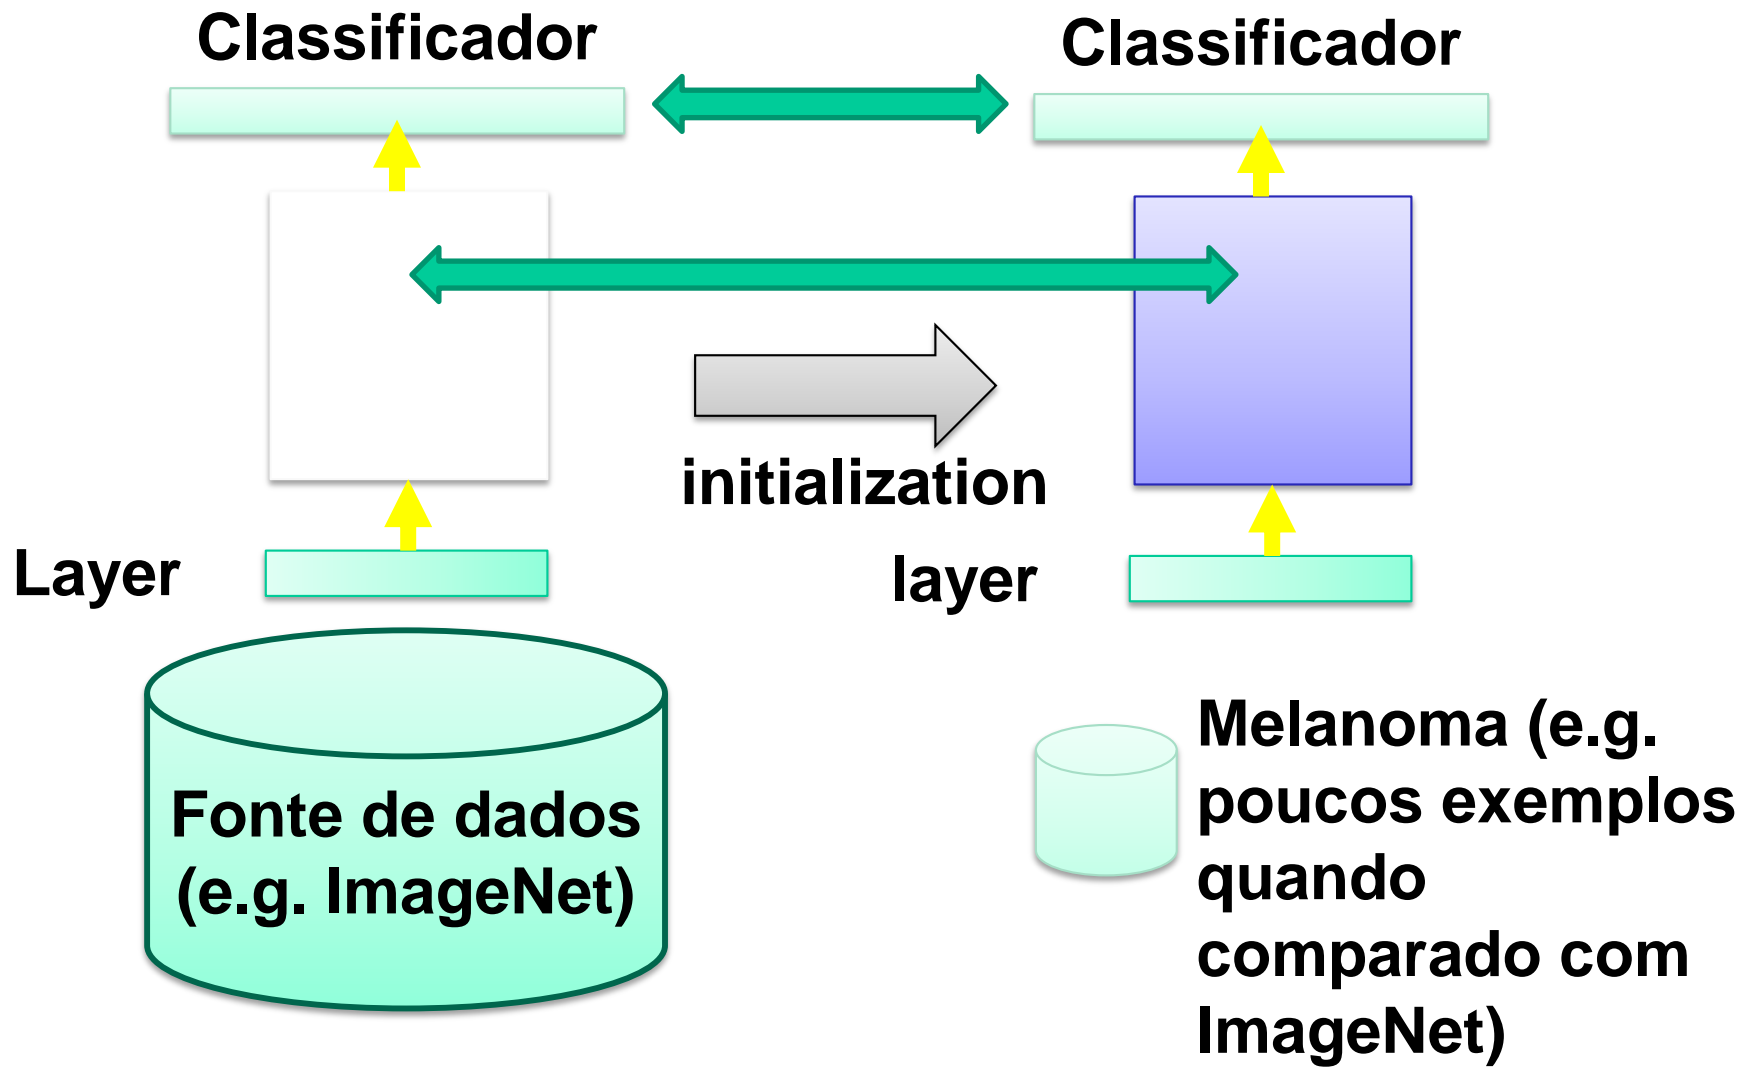
- Fonte de dados: audio e transcrição de várias falas
- Alvo: transcrição para texto

Idea: Treinar um modelo a partir de uma fonte grande de dados, depois fazer um ajuste fino para a situação em que se tem poucos dados

- Desafio: cuidado com o overfitting

# Conservative Training

---




# Layer Transfer

---

Quais camadas podem ser transferidas?

- PLN: Geralmente as últimas camadas
- Imagem: Geralmente as primeiras camadas



# Transfer Learning - Overview


---

		Source Data	
		supervisionado	Não supervisionado
Target Data	Superv.	<b>Fine-tuning</b> <b>Multitask Learning</b>	<b>Self-taught learning</b> Rajat Raina , Alexis Battle , Honglak Lee , Benjamin Packer , Andrew Y. Ng, Self-taught learning: transfer learning from unlabeled data, ICML, 2007
	Não Super.	<b>Domain- adversarial training</b> <b>Zero-shot learning</b>	<b>Self-taught Clustering</b> Wenyuan Dai, Qiang Yang, Gui-Rong Xue, Yong Yu, "Self-taught clustering". ICML 2008

# Transfer Learning

---

- Extração de características: retreinar apenas o classificador (Fully Coneceted Layer)
- Aproveitar somente a arquitetura
- Reinicializar algumas camadas



# Transfer Learning

---

- **Cenário 1 – Data set pequeno e a similaridade dos dados é muito alta**
- Use a rede pretreinada como feature extractor.
- Exemplo: Suponha que vamos usar uma rede treinada para um novo conjunto de gatos e cachorros.
- ImageNet: 1000 classes de saída
- Modificamos a últimas camadas (dense layers) e a softmax para duas categorias

# Transfer Learning

---

- **Cenário 2 – Data set pequeno e similaridade de dados baixa**
- Mantenha as camadas iniciais da rede pretreinada.
- Retreine as camadas intermediárias e finais

# Transfer Learning

---

- **Cenário 3 – Data set grande e similaridade de dados baixa**
- Neste caso.... Fuja para as montanhas...





# Transfer Learning

---

- **Cenário 4 – Data set grande e alta similaridade**
- Situação ideal: no máximo treinar o classificador no topo da rede.



---

# Bônus

# Transferência de estilo

---

## A neural Algorithm of Artistic Style

<https://arxiv.org/pdf/1508.06576.pdf>

### 1 Upload photo

The first picture defines the scene you would like to have painted.



**Entrada**

### 2 Choose style

Choose among predefined styles or upload your own style image.



**Imagem de estilo**

### 3 Submit

Our servers paint the image for you. You get an email when it's done.



# Transferência de estilo

**Perceptual Loss** for Real-Time Style transfer  
and Super-Resolution



**Per-Pixel Loss :  $y - y^{\wedge}$**

# Transferência de estilo

---

## Feature Reconstruction Loss


$$\ell_{feat}^{\phi,j}(\hat{y}, y) = \frac{1}{C_j H_j W_j} \|\phi_j(\hat{y}) - \phi_j(y)\|_2^2$$

+

## Style Reconstruction Loss

$$\ell_{style}^{\phi,j}(\hat{y}, y) = \|G_j^{\phi}(\hat{y}) - G_j^{\phi}(y)\|_F^2.$$

# Transferência de estilo



# Transferência de estilo

---

**Conceito de "reconstrução de estilo": Em vez de tentar combinar as ativações, tente combinar a correlação das ativações.**

**O recurso de correlação é chamado de matriz de Gram: produto de ponto entre a matriz de ativação de característica vetorial e sua transposição.**

---

**A definição da matriz de Gram pode parecer confusa. Tradução literal da equação 3 do artigo original usando numpy:**

```
def gram(layer):  
    N = layer.shape[1]  
    F = layer.reshape(N, -1)  
    M = F.shape[1]  
    G = np.zeros((N, N))  
    for i in range(N):  
        for j in range(N):  
            for k in range(M):  
                G[i,j] += F[i,k] * F[j,k]  
    return G
```



Style  
*The Starry Night*,  
 Vincent van Gogh,  
 1889



Style  
*The Muse*,  
 Pablo Picasso,  
 1935



Style  
*Composition VII*,  
 Wassily  
 Kandinsky, 1913



Style  
*The Great Wave off*  
*Kanagawa*, Hokusai,  
 1829-1832



# Lista de leituras

---

- <https://culurciello.github.io/tech/2016/06/04/nets.html>
- Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, [Gradient-based learning applied to document recognition](#), Proc. IEEE 86(11): 2278–2324, 1998.
- A. Krizhevsky, I. Sutskever, and G. Hinton, [ImageNet Classification with Deep Convolutional Neural Networks](#), NIPS 2012
- M. Zeiler and R. Fergus, [Visualizing and Understanding Convolutional Networks](#), ECCV 2014
- K. Simonyan and A. Zisserman, [Very Deep Convolutional Networks for Large-Scale Image Recognition](#), ICLR 2015
- M. Lin, Q. Chen, and S. Yan, [Network in network](#), ICLR 2014
- C. Szegedy et al., [Going deeper with convolutions](#), CVPR 2015
- C. Szegedy et al., [Rethinking the inception architecture for computer vision](#), CVPR 2016
- K. He, X. Zhang, S. Ren, and J. Sun, [Deep Residual Learning for Image Recognition](#), CVPR 2016

---

<https://datascience.stackexchange.com/questions/15328/what-is-the-difference-between-inception-v2-and-inception-v3>

<https://pdfs.semanticscholar.org/73ac/009051bba99eaea799172b28d69168b6aa02.pdf>

<http://blog.csdn.net/u010025211/article/details/51206237>

<http://lsun.cs.princeton.edu/slides/Christian.pdf>

<https://sourcedexter.com/retrain-tensorflow-inception-model/>

---

<https://www.quora.com/How-are-1x1-convolutions-used-for-dimensionality-reduction>

<http://www.ic.unicamp.br/~rocha/teaching/2015s2/mo826/classes/2015-google-faces.pdf>

[http://www.cv-foundation.org/openaccess/content\\_cvpr\\_2015/papers/Szegedy\\_Going\\_Deeper\\_With\\_2015\\_CVPR\\_paper.pdf](http://www.cv-foundation.org/openaccess/content_cvpr_2015/papers/Szegedy_Going_Deeper_With_2015_CVPR_paper.pdf)

[http://cs231n.stanford.edu/slides/2017/cs231n\\_2017\\_lecture9.pdf](http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture9.pdf)

[https://www.robots.ox.ac.uk/~vgg/rg/slides/rg\\_4feb16.pdf](https://www.robots.ox.ac.uk/~vgg/rg/slides/rg_4feb16.pdf)